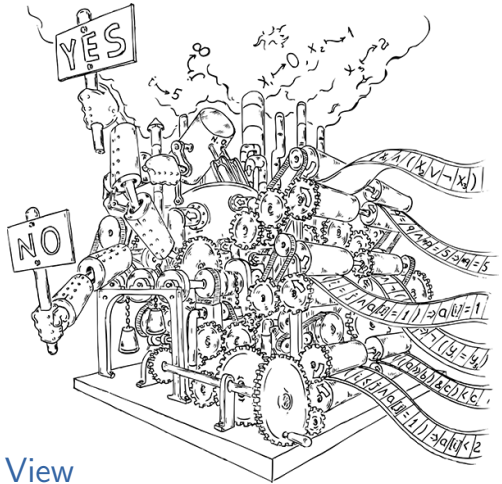


Propositional Encodings

Chapter 11



Decision Procedures An Algorithmic Point of View

- ➊ Overview
- ➋ Notation
- ➌ A Basic Encoding Algorithm
- ➍ Integration into DPLL
- ➎ DPLL(T)
- ➏ Optimizations and Implementation Issues

- Let T be a first-order Σ -theory such that:
 - T is quantifier-free.
 - There exists a decision procedure, denoted DP_T , for the conjunctive fragment of T .

- Example 1:
 - T is equality logic.
 - DP_T is an algorithm based on union-find.

- Example 2:
 - T is disjunctive linear arithmetic.
 - DP_T is the Simplex algorithm.

Example: deciding a conjunction of equalities

Input: a conjunction of equalities and disequalities φ .

Example:

$$\varphi : \quad x_1 = x_2 \wedge x_2 = x_3 \wedge x_4 = x_5 \wedge x_5 \neq x_1$$

Example: deciding a conjunction of equalities

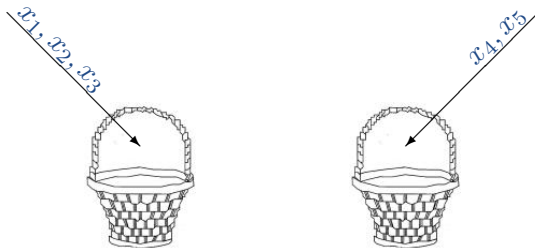
Algorithm

- ➊ Define an **equivalence class** for each variable.
- ➋ For each equality $x = y$, merge the classes of x and y .
- ➌ For each disequality $x \neq y$:
if x is in the same class as y , return 'UNSAT'.
- ➍ Return 'SAT'.

Can be implemented efficiently with a **union-find algorithm**.

Example: deciding a conjunction of equalities

$$\varphi: \quad x_1 = x_2 \wedge x_2 = x_3 \wedge x_4 = x_5 \wedge x_5 \neq x_1$$



Equivalence class 1 Equivalence class 2

Q: Is there a disequality between members of the same class?

We will now study a framework that combines

- DP_T , and
- a SAT solver,

in various ways, in order to construct a decision procedure for T .

This method is

- modular,
- efficient,
- competitive (all state-of-the-art SMT solvers work this way).

The two main engines in this framework work in tight collaboration:

- The **SAT solver** chooses those literals that need to be satisfied in order to satisfy the Boolean structure of the formula, and
- The **theory solver** DP_T checks whether this choice is consistent in T .

Let l be a Σ -literal.

- Denote by $e(l)$ the **Boolean encoder of this literal**.

Let φ be a Σ -formula,

- Denote by $e(\varphi)$ the Boolean formula resulting from substituting each Σ -literal in φ with its Boolean encoder.

For a Σ -formula φ , the resulting Boolean formula $e(\varphi)$ is called the **propositional skeleton of φ** .

- **Example I:** Let $l := x = y$ be a Σ -literal. Then $e(x = y)$, a Boolean variable, is its encoder.

- **Example II:** Let

$$\varphi := x = y \vee x = z$$

be a Σ -formula. Then

$$e(\varphi) := e(x = y) \vee e(x = z)$$

is its Boolean encoder.

Let T be equality logic. Given an NNF formula

$$\varphi := x = y \wedge ((y = z \wedge x \neq z) \vee x = z) , \quad (1)$$

we begin by computing its propositional skeleton:

$$e(\varphi) := e(x = y) \wedge ((e(y = z) \wedge e(x \neq z)) \vee e(x = z)) . \quad (2)$$

Note that since we are encoding *literals* and not *atoms*, $e(\varphi)$ has no negations and hence is trivially satisfiable.

Let \mathcal{B} be a Boolean formula, initially set to $e(\varphi)$, i.e.,

$$\mathcal{B} := e(\varphi) .$$

As a second step, we pass \mathcal{B} to a SAT solver.

Assume that the SAT solver returns the satisfying assignment

$$\alpha := \{e(x = y) \mapsto \text{TRUE}, e(y = z) \mapsto \text{TRUE}, e(x \neq z) \mapsto \text{TRUE}, \\ e(x = z) \mapsto \text{FALSE}\} .$$

- Denote by $\hat{T}h(\alpha)$ the conjunction of the literals corresponding to this assignment.

$$\hat{T}h(\alpha) := x = y \wedge y = z \wedge x \neq z \wedge \neg(x = z) .$$

- The decision procedure DP_T now has to decide whether $\hat{T}h(\alpha)$ is satisfiable.

$\hat{T}h(\alpha)$ is not satisfiable, which means that the negation of this formula is a tautology.

Thus \mathcal{B} is conjoined with $e(\neg\hat{T}h(\alpha))$, the Boolean encoding of this tautology:

$$e(\neg\hat{T}h(\alpha)) := (\neg e(x = y) \vee \neg e(y = z) \vee \neg e(x \neq z) \vee e(x = z)) .$$

- This clause contradicts the current assignment, and hence **blocks** it from being repeated.
- Such clauses are called **blocking clauses**.
- We denote by t the formula – also called the **lemma** – returned by DP_T (in this example $t := \neg\hat{T}h(\alpha)$).

After the blocking clause has been added, the SAT solver is invoked again and suggests another assignment, for example

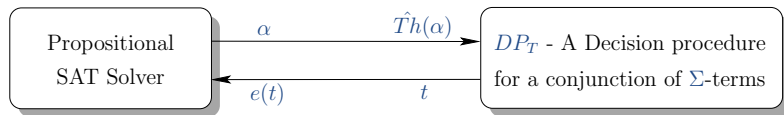
$$\alpha' := \{e(x = y) \mapsto \text{TRUE}, e(y = z) \mapsto \text{TRUE}, e(x = z) \mapsto \text{TRUE}, \\ e(x \neq z) \mapsto \text{FALSE}\}.$$

The corresponding Σ -formula

$$\hat{T}h(\alpha') := x = y \wedge y = z \wedge x = z \wedge \neg(x \neq z) \quad (3)$$

is satisfiable, which proves that φ , the original formula, is satisfiable.

Indeed, any assignment that satisfies $\hat{T}h(\alpha')$ also satisfies φ .



The information flow between the two components of the decision procedure.

There are many improvements to this basic procedure.

We will later consider several of them:

- ❶ Invoke DP_T after **partial assignments**.
- ❷ **Theory propagation**: learn propositional constraints based on the theory literals.
 - When the partial assignment is not contradictory.
- ❸ **Generalize** blocking clauses

Consider the partial assignment

$$\alpha := \{e(x = y) \mapsto \text{TRUE}, e(y = z) \mapsto \text{TRUE}\} . \quad (4)$$

DP_T receives:

$$\hat{Th}(\alpha) := x = y \wedge y = z , \quad (5)$$

... and infers $x = z$.

DP_T now performs **Theory Propagation**: informs the SAT solver of

$$e(x = z) \mapsto \text{TRUE} \text{ and } e(x \neq z) \mapsto \text{FALSE}.$$

We will now formalize three versions of the algorithm:

- ➊ Simple
- ➋ Incremental
- ➌ DPLL(T)

- $lit(\varphi)$ – the set of literals in a given NNF formula φ .
- $lit_i(\varphi)$ – the i -th distinct literal in φ
(assuming some predefined order on the literals).
- α – For a given encoding $e(\varphi)$, denotes an assignment (either full or partial), to the encoders in $e(\varphi)$.

- $Th(lit_i, \alpha)$ – For an encoder $e(lit_i)$ that is assigned a truth value by α , denotes the corresponding literal:

$$Th(lit_i, \alpha) \doteq \begin{cases} lit_i & \alpha(lit_i) = \text{TRUE} \\ \neg lit_i & \alpha(lit_i) = \text{FALSE} . \end{cases} \quad (6)$$

- $Th(\alpha) \doteq \{Th(lit_i, \alpha) \mid e(lit_i) \text{ is assigned by } \alpha\}$
- $\hat{Th}(\alpha)$ – a conjunction over the elements in $Th(\alpha)$.

Let

$$lit_1 = (x = y), \quad lit_2 = (y = z), \quad lit_3 = (z = w) , \quad (7)$$

and let α be a partial assignment such that

$$\alpha := \{e(lit_1) \mapsto \text{FALSE}, \quad e(lit_2) \mapsto \text{TRUE}\} .$$

Then

$$Th(lit_1, \alpha) := \neg(x = y), \quad Th(lit_2, \alpha) := (y = z) ,$$

and

$$Th(\alpha) := \{\neg(x = y), (y = z)\} .$$

Conjoining these terms gives us

$$\hat{Th}(\alpha) := \neg(x = y) \wedge (y = z) .$$

- T – a Σ -theory.
- DP_T a decision procedure for the conjunctive fragment of T .
- Let DEDUCTION be a procedure based on DP_T , which receives a conjunction of Σ -literals as input, and
 - decides whether it is satisfiable, and,
 - if the answer is negative, returns constraints over these literals.

1. A Basic Algorithm

```
1: function LAZY-BASIC( $\varphi$ )
2:    $\mathcal{B} := e(\varphi)$ ;
3:   while (TRUE) do
4:      $\langle \alpha, res \rangle := \text{SAT-SOLVER}(\mathcal{B})$ ;
5:     if  $res = \text{"Unsatisfiable"}$  then return "Unsatisfiable";
6:     else
7:        $\langle t, res \rangle := \text{DEDUCTION}(\hat{T}h(\alpha))$ ;
8:       if  $res = \text{"Satisfiable"}$  then return "Satisfiable";
9:        $\mathcal{B} := \mathcal{B} \wedge e(t)$ ;
```

1. DEDUCTION

The clause t that is returned by DEDUCTION:

... should not be **too strong**.

Strategies:

- ❶ $\varphi \rightarrow t$.
 - ... DEDUCTION needs to deal with φ rather than $Atoms(\varphi)$.
- ❷ t is T -valid

Example:

$$x = y \wedge y = z \longrightarrow x = z$$

This guarantees **soundness**.

1. DEDUCTION

The clause t that is returned by DEDUCTION:

... should not be too weak

- it should at least block α .

This guarantees termination (?)



1. DEDUCTION

The clause t that is returned by DEDUCTION:

... should not lead to **divergence**

- $Atoms(t) \subseteq Atoms(\varphi)$, or
- $Atoms(t)$ is finite.

Example:

T = equality logic.

$$Atoms(t) = \{x_i = x_j \mid x_i, x_j \in var(\varphi)\}$$

Some of the predicates in t do not appear in φ .

This guarantees **termination**.

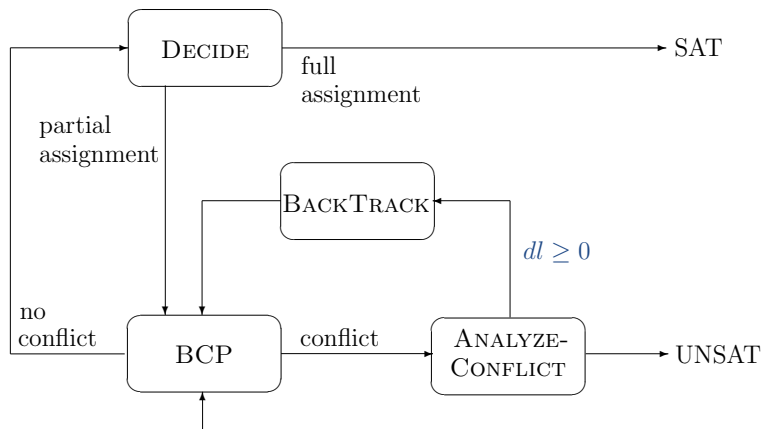
2. We can do better...

- Let \mathcal{B}^i be the formula \mathcal{B} in the i -th iteration of the loop.
- The constraint \mathcal{B}^{i+1} is strictly stronger than \mathcal{B}^i for all $i \geq 1$, because clauses are added but not removed between iterations.
- As a result, any conflict clause that is learned while solving \mathcal{B}^i can be reused when solving \mathcal{B}^j for $i < j$.
- This is a special case of **incremental satisfiability**.

2. We can do better...

- Hence, invoking an incremental SAT solver in line 4 can increase the efficiency of the algorithm.
- A better option is to **integrate** DEDUCTION into the DPLL-SAT algorithm, as shown in the following algorithm.
- This algorithm uses a procedure **ADDCLAUSES**, which adds new clauses to the current set of clauses at run time.
- Before seeing this algorithm let us first recall DPLL...

2. A Reminder: DPLL



2. Pseudo-code for DPLL

```
1: function DPLL
2:   if BCP() = "conflict" then return "Unsatisfiable";
3:   while (TRUE) do
4:     if  $\neg$ DECIDE() then return "Satisfiable";
5:     while (BCP() = "conflict") do
6:       bcktk-level := ANALYZE-CONFLICT();
7:       if bcktk-level < 0 then return "Unsatisfiable";
8:       else BackTrack(bcktk-level);
```


2. Integration into DPLL

```
1: function LAZY-DPLL
2:   ADDCLAUSES(cnf(e( $\varphi$ )));
3:   if BCP() = "conflict" then return "Unsatisfiable";
4:   while (TRUE) do
5:     if  $\neg$ DECIDE() then ▷ Full assignment
6:        $\langle t, res \rangle :=$  DEDUCTION( $\hat{T}h(\alpha)$ );
7:       if res = "Satisfiable" then return "Satisfiable";
8:       ADDCLAUSES(e(t));
9:     while (BCP() = "conflict") do
10:      backtrack-level := ANALYZE-CONFLICT();
11:      if backtrack-level < 0 then return "Unsatisfiable";
12:      else BackTrack(backtrack-level);
```

3. ... or even better

- Let $\{x_1 \geq 10, x_1 < 0\} \subset Atoms(\varphi)$.
- Assume currently $e(x_1 \geq 10) = e(x_1 < 0) = \text{TRUE}$.
- Now any call to DEDUCTION results in a **contradiction**.

3. ... or even better

- So far: DEDUCTION after finding a full assignment α .
 - Time taken to complete the assignment is wasted.
 - The refutation of α may be due to other reasons.
 - additional assignments that include $e(x_1 \geq 10) = e(x_1 < 0) = \text{TRUE}$ are not ruled out.

Solution: Early call to Deduction.

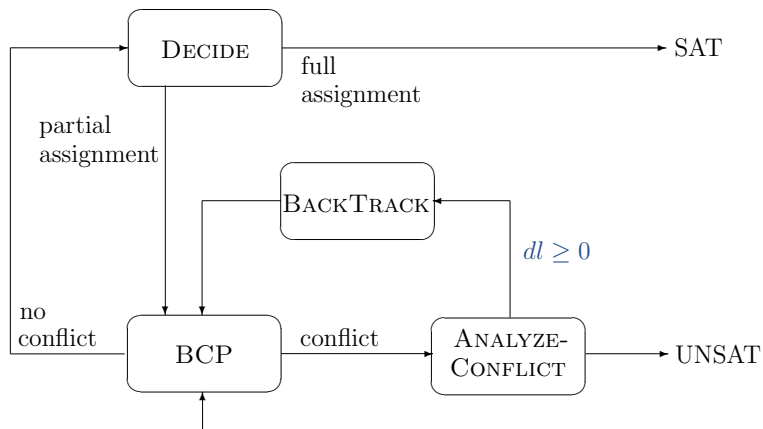
3. The DPLL(T) Framework

Early call to DEDUCTION can serve two purposes:

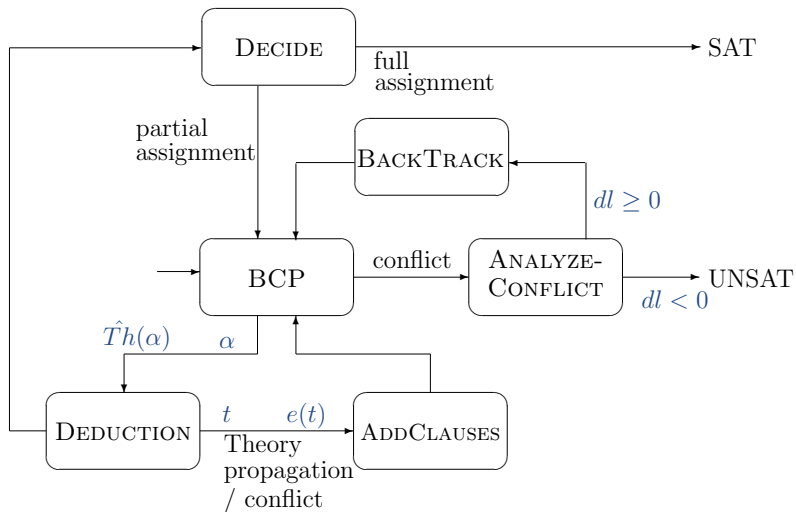
- ❶ Contradictory partial assignments are ruled out early.
- ❷ Allows **theory propagation**.
 - When $e(x_1 \geq 10) = \text{TRUE}$, infer $e(x_1 < 0) = \text{FALSE}$.

This brings us to the next version of the algorithm, called DPLL(T).

3. Reminder: DPLL



3. ... and now DPLL(T)



```

1: function DPLL( $T$ )
2:   ADDCLAUSES( $cnf(e(\varphi))$ );
3:   if BCP() = "conflict" then return "Unsatisfiable";
4:   while (TRUE) do
5:     if  $\neg$ DECIDE() then return "Satisfiable";
6:     repeat
7:       while (BCP() = "conflict") do
8:          $btrk\text{-}level :=$  ANALYZE-CONFLICT();
9:         if  $btrk\text{-}level < 0$  then return "Unsatisfiable";
10:        else BackTrack( $btrk\text{-}level$ );
11:         $\langle t, res \rangle :=$  DEDUCTION( $\hat{T}h(\alpha)$ );
12:        ADDCLAUSES( $e(t)$ );
13:    until  $t \equiv \text{TRUE}$ 

```

3. Restrictions on DEDUCTION

When $\hat{T}h(\alpha)$ is satisfiable, it is required that:

- $e(t)$ is an **asserting clause** under α .
 - BCP will now assign a value to an encoder.
- If DEDUCTION cannot find such an asserting clause t , then $t = e(t) = \text{TRUE}$.
 - This can happen when, e.g., all the encoders are assigned.

3. Theory Propagation

Various ways to perform theory propagation:

- After every **decision** / after every **assignment**
- **Partial** / **Exhaustive** theory propagation
 - Exhaustive = propagate everything that is implied.
- Refer only to **existing** predicates / **add** auxiliary ones.

Exhaustive theory propagation after each assignment: what does this mean ?

That's right, no possible conflicts on the theory side.

3. Theory Propagation

How to check whether a predicate p is implied by $\hat{Th}(\alpha)$?

- **Plunging** – is $\hat{Th}(\alpha) \wedge \neg p$ satisfiable ?
- Theory-specific propagation.

Example:

T = equality logic

Build the equality graph corresponding to $Th(\alpha)$.

Infer equalities/disequalities from the graph.

3. Theory Propagation: observations

- Theory propagation matters for efficiency, not correctness.
- How much propagation is cost-effective is a subject for research, and depends on T .

3. Theory Propagation – How?

- Normally theory propagation is done by transferring clauses.
- Inefficient
 - Less than 0.5% are actually used.
- Instead – add implied literals directly to the implication stack.
 - This causes a **problem** in `ANALYZE-CONFLICT()`
 - Can you see what problem ?

3. Theory Propagation – How?

- The problem: $\text{ANALYZE-CONFLICT}()$ requires an **antecedent clause** for each implication, in order to compute the conflict clause and backtrack level.
- ... but now there are implications without antecedents.
- Solution – DP_T should be able to **explain** an implication post-mortem, in the form of a clause.

3. Strong Lemmas

- When $\hat{T}h(\alpha)$ is unsatisfiable, the ‘lemma’ t blocks α .
- We want to make t stronger.
- How?
 - Analyze the **reason** for the unsatisfiability.
 - Build the lemma accordingly.

3. Strong Lemmas – An Example

